

Hidden Markov Models

Hrant Baloyan, Khachatur Inants

Hidden Markov Models (HMMs) are an extension of Markov Chains.

Markov Chain consists of

1) A sequence of states


2) A transition matrix

3) An initial probability distribution

Main property

$$\textbf{Markov Assumption: } P(q_i = a | q_1 \dots q_{i-1}) = P(q_i = a | q_{i-1})$$

Example of Markov Chain

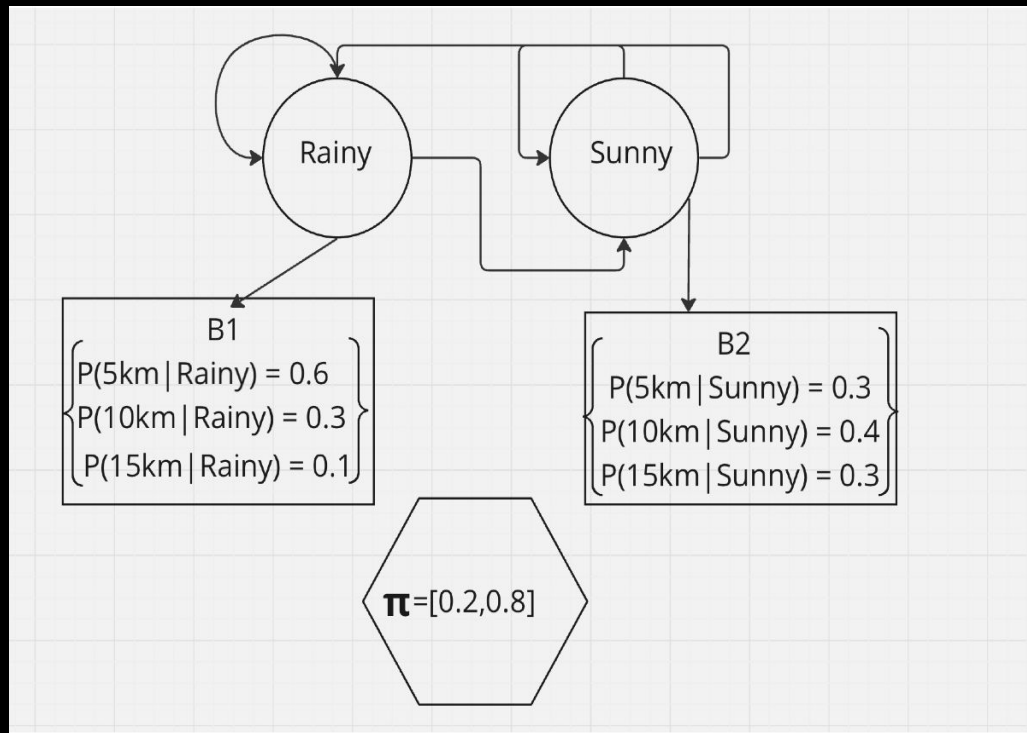


	clouds	rain	sun
clouds	0.4	0.3	0.3
rain	0.5	0.3	0.2
sun	0.5	0.1	0.4

But what if we have hidden events?

HMM has

- 1) Sequence of hidden states
- 2) Transition matrix
- 3) Probabilities that represent what's the probability of observation o being generated by state q
- 4) An Initial probability distribution



HMM makes two assumptions

Markov assumption:

$$P(q_i | q_1 \cdots q_{i-1}) = P(q_i | q_{i-1})$$

Independence assumption:

$$P(o_i | q_1 \cdots q_i \cdots q_T, o_1, \cdots o_i, \cdots o_T) = P(o_i | q_i)$$

HMMs can be described in 3 problems.

- 1) Likelihood - Given HMM $\lambda = (A, B)$ and sequence of observations O , find $P(O|\lambda)$
- 1) Decoding - Given sequence of observations O and HMM $\lambda = (A, B)$, find best sequence of hidden states
- 1) Likelihood - Given HMM $\lambda = (A, B)$ and sequence of observations O , find $P(O|\lambda)$

Likelihood Computation : The Forward Algorithm

Problem is to compute the likelihood of a particular observation sequence.

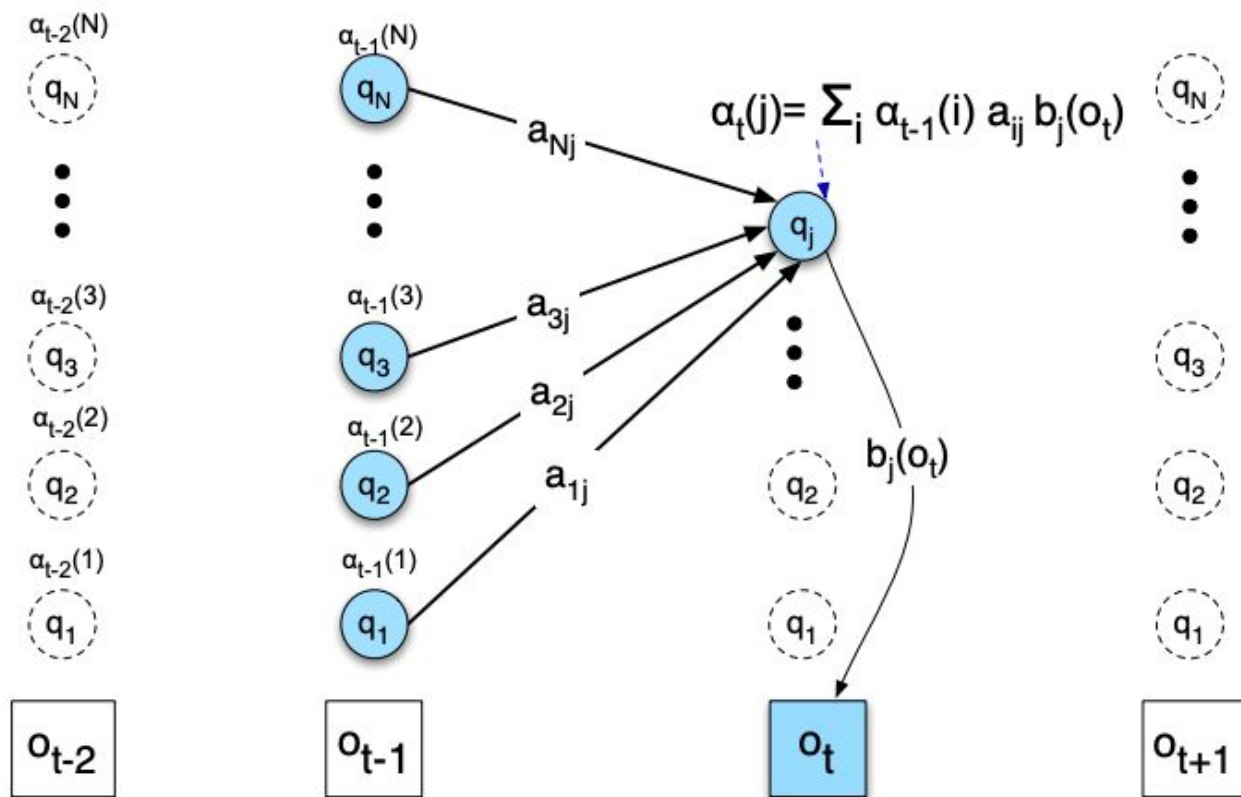
Each cell of the forward algorithm trellis $\alpha_t(j)$ represents the probability of being in state j after seeing the first t observations,

$$\alpha_t(j) = \sum_{i=1}^N \alpha_{t-1}(i) a_{ij} b_j(o_t)$$

$\alpha_{t-1}(i)$ the previous forward path probability from the previous time step

a_{ij} the transition probability from previous state q_i to current state q_j

$b_j(o_t)$ the state observation likelihood of the observation symbol o_t given the current state j



function FORWARD(*observations* of len T , *state-graph* of len N) **returns** *forward-prob*

create a probability matrix *forward*[N, T]

for each state s **from** 1 **to** N **do** ; initialization step

$forward[s, 1] \leftarrow \pi_s * b_s(o_1)$

for each time step t **from** 2 **to** T **do** ; recursion step

for each state s **from** 1 **to** N **do**

$$forward[s, t] \leftarrow \sum_{s'=1}^N forward[s', t-1] * a_{s', s} * b_s(o_t)$$

$$forwardprob \leftarrow \sum_{s=1}^N forward[s, T] \quad ; \text{termination step}$$

return *forwardprob*

The Viterbi algorithm

Given a series of observed events, the Viterbi algorithm determines the most likely order of hidden states in an HMM

$v_t(j)$, represents the probability that the HMM is in state j after seeing the first t observations and passing through the most probable state sequence q_1, \dots, q_{t-1} ,

$$v_t(j) = \max_{i=1}^N v_{t-1}(i) a_{ij} b_j(o_t)$$

$v_{t-1}(i)$ the previous Viterbi path probability from the previous time step

a_{ij} the transition probability from previous state q_i to current state q_j

$b_j(o_t)$ the state observation likelihood of the observation symbol o_t given the current state j

$bt_t(j)$, represents which state at a time $t-1$ maximize the probability of reaching the state j after seeing the observation t .

$$bt_t(j) = \underset{i=1}{\operatorname{argmax}}^N v_{t-1}(i) a_{ij} b_j(o_t);$$

```

for each state  $s$  from 1 to  $N$  do                                ; initialization step
     $viterbi[s,1] \leftarrow \pi_s * b_s(o_1)$ 
     $backpointer[s,1] \leftarrow 0$ 
for each time step  $t$  from 2 to  $T$  do                                ; recursion step
    for each state  $s$  from 1 to  $N$  do
         $viterbi[s,t] \leftarrow \max_{s'=1}^N viterbi[s',t-1] * a_{s',s} * b_s(o_t)$ 
         $backpointer[s,t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s',t-1] * a_{s',s} * b_s(o_t)$ 

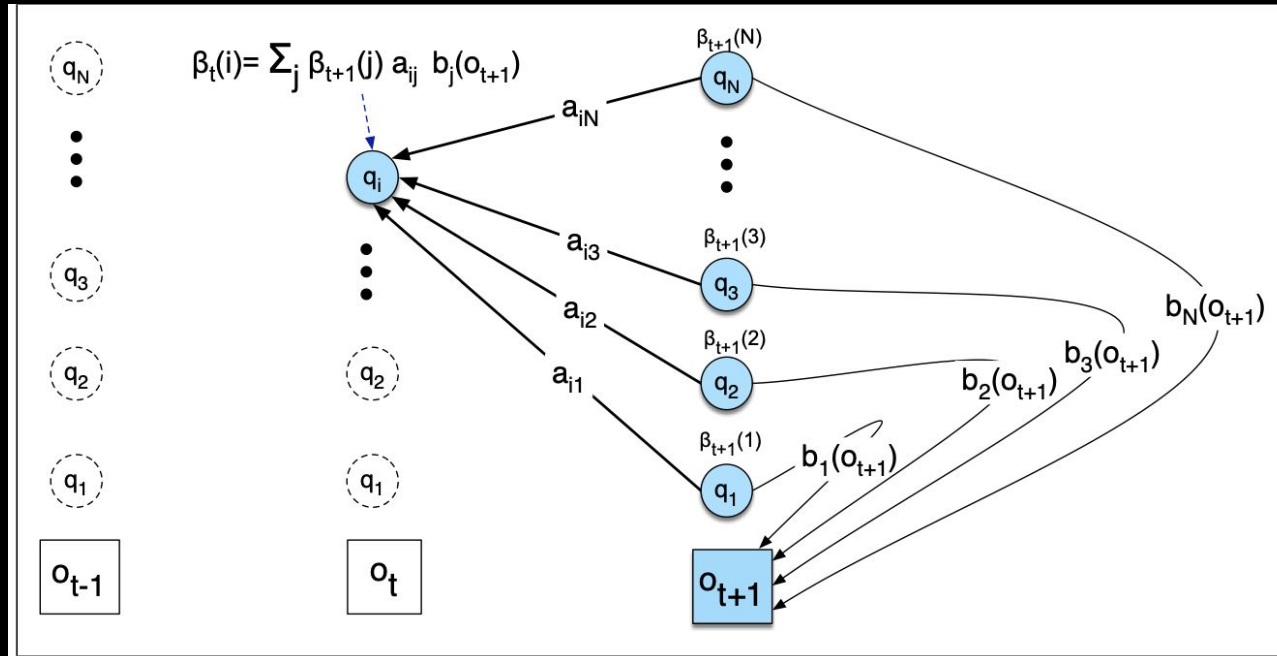
     $bestpathprob \leftarrow \max_{s=1}^N viterbi[s,T]$                                 ; termination step
     $bestpathpointer \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s,T]$                                 ; termination step
     $bestpath \leftarrow$  the path starting at state  $bestpathpointer$ , that follows  $backpointer[]$  to states back in time
return  $bestpath, bestpathprob$ 

```

Learning of HMM: Forward-Backward algorithm

Here we are trying to find parameters of HMM A and B using observations and states

Introducing backward probabilities



Need we try to estimate a_{ij}

$$\hat{a}_{ij} = \frac{\text{expected number of transitions from state } i \text{ to state } j}{\text{expected number of transitions from state } i}$$

Let's define


$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j | O, \lambda)$$

First let's compute

$$\text{not-quite-}\xi_t(i, j) = P(q_t = i, q_{t+1} = j, O | \lambda)$$

Continue

$$\text{not-quite-}\xi_t(i, j) = \alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)$$



Using

$$P(X|Y,Z) = \frac{P(X,Y|Z)}{P(Y|Z)}$$

We compute ξ_t from not-quite- ξ_t , by dividing on $P(O|\lambda)$

$$P(O|\lambda) = \sum_{j=1}^N \alpha_t(j) \beta_t(j)$$

Then final equation for ξ_t is

$$\xi_t(i,j) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\sum_{j=1}^N \alpha_t(j) \beta_t(j)}$$

Then expected number of transitions from state i to state j is then the sum over all t of ξ

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i,j)}{\sum_{t=1}^{T-1} \sum_{k=1}^N \xi_t(i,k)}$$

And now try to estimate probability of a given symbol v_k from the observation vocabulary V , given a state j

$$\hat{b}_j(v_k) = \frac{\text{expected number of times in state } j \text{ and observing symbol } v_k}{\text{expected number of times in state } j}$$

Calculate For this we need to compute

$$\gamma_t(j) = P(q_t = j | O, \lambda)$$

Same way

$$\gamma_t(j) = \frac{P(q_t = j, O | \lambda)}{P(O | \lambda)}$$

$$\gamma_t(j) = \frac{\alpha_t(j) \beta_t(j)}{P(O | \lambda)}$$

$$\hat{b}_j(v_k) = \frac{\sum_{t=1}^T \mathbb{1}_{O_t=v_k} \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)}$$

function FORWARD-BACKWARD(*observations* of len T , *output vocabulary* V , *hidden state set* Q) **returns** $HMM=(A,B)$

initialize A and B

iterate until convergence

E-step

$$\gamma_t(j) = \frac{\alpha_t(j)\beta_t(j)}{\alpha_T(q_F)} \quad \forall t \text{ and } j$$

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\alpha_T(q_F)} \quad \forall t, i, \text{ and } j$$

M-step

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \sum_{k=1}^N \xi_t(i, k)}$$

$$\hat{b}_j(v_k) = \frac{\sum_{t=1}^T \mathbb{1}_{O_t=v_k} \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)}$$

return A, B



```
for i in range(10):  
    print(train_data[i][:7])
```



```
[('Pierre', 'NOUN'), ('Vinken', 'NOUN'), (',', '.'), ('61', 'NUM'), ('years', 'NOUN'), ('old', 'ADJ'), (',', '.')]  
[('Mr.', 'NOUN'), ('Vinken', 'NOUN'), ('is', 'VERB'), ('chairman', 'NOUN'), ('of', 'ADP'), ('Elsevier', 'NOUN'), ('N.V.', 'NOUN')]  
[('Rudolph', 'NOUN'), ('Agnew', 'NOUN'), (',', '.'), ('55', 'NUM'), ('years', 'NOUN'), ('old', 'ADJ'), ('and', 'CONJ')]  
[('A', 'DET'), ('form', 'NOUN'), ('of', 'ADP'), ('asbestos', 'NOUN'), ('once', 'ADV'), ('used', 'VERB'), ('*', 'X')]  
[('The', 'DET'), ('asbestos', 'NOUN'), ('fiber', 'NOUN'), (',', '.'), ('crocidolite', 'NOUN'), (',', '.'), ('is', 'VERB')]  
[('Lorillard', 'NOUN'), ('Inc.', 'NOUN'), (',', '.'), ('the', 'DET'), ('unit', 'NOUN'), ('of', 'ADP'), ('New', 'ADJ')]  
[('Although', 'ADP'), ('preliminary', 'ADJ'), ('findings', 'NOUN'), ('were', 'VERB'), ('reported', 'VERB'), ('*-2', 'X'), ('more', 'ADV')]  
[('A', 'DET'), ('Lorillard', 'NOUN'), ('spokewoman', 'NOUN'), ('said', 'VERB'), (',', '.'), ('`', '.'), ('This', 'DET')]  
[('We', 'PRON'), ('re', 'VERB'), ('talking', 'VERB'), ('about', 'ADP'), ('years', 'NOUN'), ('ago', 'ADP'), ('before', 'ADP')]  
[('There', 'DET'), ('is', 'VERB'), ('no', 'DET'), ('asbestos', 'NOUN'), ('in', 'ADP'), ('our', 'PRON'), ('products', 'NOUN')]
```

Test Sentence: ['the', 'quick', 'brown', 'fox', 'jumps', 'over', 'the', 'lazy', 'dog.']

Predicted Tags: ['DET', 'ADJ', 'NOUN', 'ADP', 'NOUN', 'ADP', 'DET', 'ADJ', 'NOUN']

Accuracy: 88.91%